

```
In [101]: # Jump-Start Example: Python analysis of MSPA Software Survey

# Update 2017-09-21 by Tom Miller and Kelsey O'Neill
# Update 2018-06-30 by Tom Miller v005 transformation code added

# tested under Python 3.6.1 :: Anaconda custom (x86_64)
# on Windows 10.0 and Mac OS Sierra 10.12.2

# shows how to read in data from a comma-delimited text file
# manipulate data, create new count variables, define categorical variables,
# work with dictionaries and lambda mapping functions for recoding data

# visualizations in this program are routed to external pdf files
# so they may be included in printed or electronic reports

# prepare for Python version 3x features and functions
# these two lines of code are needed for Python 2.7 only
# commented out for Python 3.x versions
# from __future__ import division, print_function
# from future_builtins import ascii, filter, hex, map, oct, zip

# external libraries for visualizations and data manipulation
# ensure that these packages have been installed prior to calls
import pandas as pd # data frame operations
import numpy as np # arrays and math functions
import matplotlib.pyplot as plt # static plotting
import seaborn as sns # pretty plotting, including heat map
```

```
In [102]: # correlation heat map setup for seaborn
def corr_chart(df_corr):
    corr=df_corr.corr()
    #screen top half to get a triangle
    top = np.zeros_like(corr, dtype=np.bool)
    top[np.triu_indices_from(top)] = True
    fig=plt.figure()
    fig, ax = plt.subplots(figsize=(12,12))
    sns.heatmap(corr, mask=top, cmap='coolwarm',
                center = 0, square=True,
                linewidths=.5, cbar_kws={'shrink':.5},
                annot = True, annot_kws={'size': 9}, fmt = '.3f')
    plt.xticks(rotation=45) # rotate variable labels on columns (x axis)
    plt.yticks(rotation=0) # use horizontal variable labels on rows (y axis)
    plt.title('Correlation Heat Map')
    plt.savefig('plot-corr-map.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)

np.set_printoptions(precision=3)
```

```
In [103]: # read in comma-delimited text file, creating a pandas DataFrame object
# note that IPAddress is formatted as an actual IP address
# but is actually a random-hash of the original IP address
valid_survey_input = pd.read_csv('mspa-survey-data.csv')
```

```
In [104]: valid_survey_input.head()
```

```
Out[104]:
```

| | RespondentID | Personal_JavaScalaSpark | Personal_JavaScriptHTMLCSS | Personal_Python | Personal_R | Persc |
|---|--------------|-------------------------|----------------------------|-----------------|------------|-------|
| 0 | 5135740122 | 0 | 0 | 0 | 50 | |
| 1 | 5133300037 | 10 | 10 | 50 | 30 | |
| 2 | 5132253300 | 20 | 0 | 40 | 40 | |
| 3 | 5132096630 | 10 | 10 | 25 | 35 | |
| 4 | 5131990362 | 20 | 0 | 0 | 70 | |

5 rows × 41 columns

```
In [105]: # use the RespondentID as label for the rows... the index of DataFrame
valid_survey_input.set_index('RespondentID', drop = True, inplace = True)
```

```
In [106]: # examine the structure of the DataFrame object
print('\nContents of initial survey data -----')
```

Contents of initial survey data -----

```
In [107]: # could use len() or first index of shape() to get number of rows/observations
print('\nNumber of Respondents =', len(valid_survey_input))
```

Number of Respondents = 207

```
In [108]: # show the column/variable names of the DataFrame
# note that RespondentID is no longer present
print(valid_survey_input.columns)
```

```
Index(['Personal_JavaScalaSpark', 'Personal_JavaScriptHTMLCSS',
      'Personal_Python', 'Personal_R', 'Personal_SAS',
      'Professional_JavaScalaSpark', 'Professional_JavaScriptHTMLCSS',
      'Professional_Python', 'Professional_R', 'Professional_SAS',
      'Industry_JavaScalaSpark', 'Industry_JavaScriptHTMLCSS',
      'Industry_Python', 'Industry_R', 'Industry_SAS',
      'Python_Course_Interest', 'Foundations_DE_Course_Interest',
      'Analytics_App_Course_Interest', 'Systems_Analysis_Course_Interest',
      'Courses_Completed', 'PREDICT400', 'PREDICT401', 'PREDICT410',
      'PREDICT411', 'PREDICT413', 'PREDICT420', 'PREDICT422', 'PREDICT450',
      'PREDICT451', 'PREDICT452', 'PREDICT453', 'PREDICT454', 'PREDICT455',
      'PREDICT456', 'PREDICT457', 'OtherPython', 'OtherR', 'OtherSAS',
      'Other', 'Graduate_Date'],
      dtype='object')
```

```
In [109]: # abbreviated printing of the first five rows of the data frame
print(pd.DataFrame.head(valid_survey_input))
```

| | Personal_JavaScalaSpark | Personal_JavaScriptHTMLCSS | \ |
|--------------|-------------------------|----------------------------|---|
| RespondentID | | | |
| 5135740122 | 0 | 0 | |
| 5133300037 | 10 | 10 | |
| 5132253300 | 20 | 0 | |
| 5132096630 | 10 | 10 | |
| 5131990362 | 20 | 0 | |

| | Personal_Python | Personal_R | Personal_SAS | \ |
|--------------|-----------------|------------|--------------|---|
| RespondentID | | | | |
| 5135740122 | 0 | 50 | 50 | |
| 5133300037 | 50 | 30 | 0 | |
| 5132253300 | 40 | 40 | 0 | |
| 5132096630 | 25 | 35 | 20 | |
| 5131990362 | 0 | 70 | 10 | |

| | Professional_JavaScalaSpark | Professional_JavaScriptHTMLCSS | \ |
|--------------|-----------------------------|--------------------------------|---|
| RespondentID | | | |
| 5135740122 | 0 | 0 | |
| 5133300037 | 25 | 25 | |
| 5132253300 | 0 | 0 | |
| 5132096630 | 10 | 10 | |
| 5131990362 | 20 | 0 | |

| | Professional_Python | Professional_R | Professional_SAS | \ |
|--------------|---------------------|----------------|------------------|---|
| RespondentID | | | | |
| 5135740122 | 0 | 25 | 75 | |
| 5133300037 | 30 | 20 | 0 | |
| 5132253300 | 40 | 40 | 20 | |
| 5132096630 | 25 | 35 | 20 | |
| 5131990362 | 0 | 80 | 0 | |

| | ... | PREDICT453 | PREDICT454 | PREDICT455 | PREDICT456 | \ |
|--------------|-----|------------|------------|------------|------------|---|
| RespondentID | ... | | | | | |
| 5135740122 | ... | NaN | NaN | NaN | NaN | |
| 5133300037 | ... | NaN | NaN | NaN | NaN | |
| 5132253300 | ... | NaN | NaN | NaN | NaN | |
| 5132096630 | ... | NaN | NaN | NaN | NaN | |
| 5131990362 | ... | NaN | NaN | NaN | NaN | |

| | PREDICT457 | OtherPython | OtherR | OtherSAS | Other | \ |
|--------------|------------|-------------|--------|----------|-------|---|
| RespondentID | | | | | | |
| 5135740122 | NaN | NaN | NaN | NaN | NaN | |
| 5133300037 | NaN | NaN | NaN | NaN | NaN | |
| 5132253300 | NaN | NaN | NaN | NaN | NaN | |
| 5132096630 | NaN | NaN | NaN | NaN | NaN | |
| 5131990362 | NaN | NaN | NaN | NaN | NaN | |

| | Graduate_Date |
|--------------|---------------|
| RespondentID | |
| 5135740122 | NaN |
| 5133300037 | Spring 2018 |
| 5132253300 | Fall 2018 |
| 5132096630 | Fall 2017 |
| 5131990362 | Fall 2018 |

[5 rows x 40 columns]

```
In [110]: # shorten the variable/column names for software preference variables
survey_df = valid_survey_input.rename(index=str, columns={
    'Personal_JavaScalaSpark': 'My_Java',
    'Personal_JavaScriptHTMLCSS': 'My_JS',
    'Personal_Python': 'My_Python',
    'Personal_R': 'My_R',
    'Personal_SAS': 'My_SAS',
    'Professional_JavaScalaSpark': 'Prof_Java',
    'Professional_JavaScriptHTMLCSS': 'Prof_JS',
    'Professional_Python': 'Prof_Python',
    'Professional_R': 'Prof_R',
    'Professional_SAS': 'Prof_SAS',
    'Industry_JavaScalaSpark': 'Ind_Java',
    'Industry_JavaScriptHTMLCSS': 'Ind_JS',
    'Industry_Python': 'Ind_Python',
    'Industry_R': 'Ind_R',
    'Industry_SAS': 'Ind_SAS'})
```

```
In [117]: # define subset DataFrame for analysis of software preferences
software_df = survey_df.loc[:, 'My_Java':'Ind_SAS']
software_df.head()
```

Out[117]:

| | My_Java | My_JS | My_Python | My_R | My_SAS | Prof_Java | Prof_JS | Prof_Python | Prof_R | Prof_SAS |
|--------------|---------|-------|-----------|------|--------|-----------|---------|-------------|--------|----------|
| RespondentID | | | | | | | | | | |
| 5135740122 | 0 | 0 | 0 | 50 | 50 | 0 | 0 | 0 | 25 | |
| 5133300037 | 10 | 10 | 50 | 30 | 0 | 25 | 25 | 30 | 20 | |
| 5132253300 | 20 | 0 | 40 | 40 | 0 | 0 | 0 | 40 | 40 | |
| 5132096630 | 10 | 10 | 25 | 35 | 20 | 10 | 10 | 25 | 35 | |
| 5131990362 | 20 | 0 | 0 | 70 | 10 | 20 | 0 | 0 | 80 | |


In [112]:

```

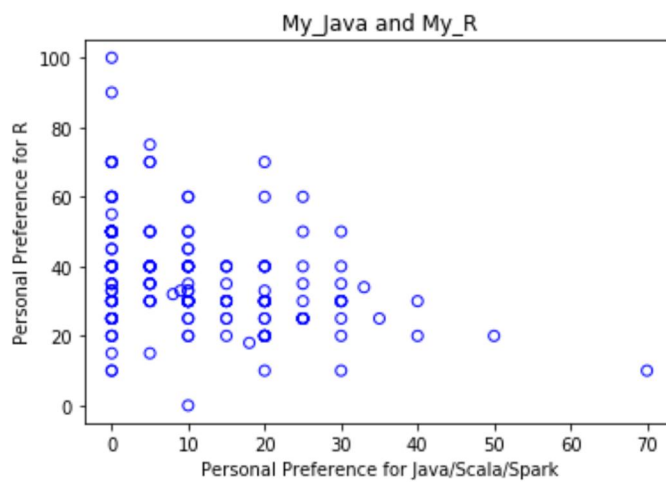

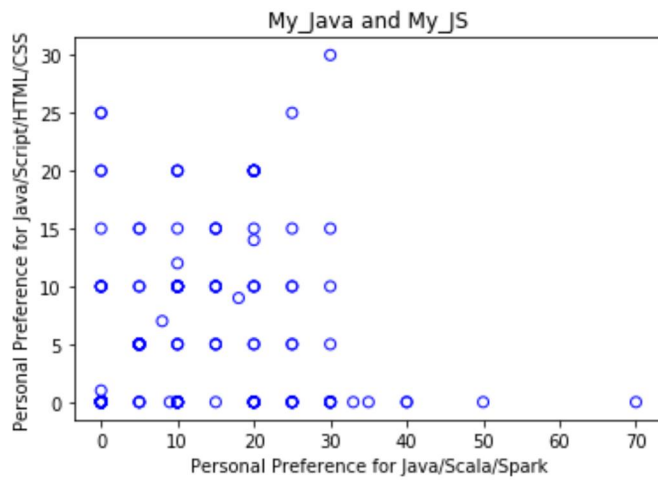
# single scatter plot example
fig, axis = plt.subplots()
axis.set_xlabel('Personal Preference for R')
axis.set_ylabel('Personal Preference for Python')
plt.title('R and Python Preferences')
scatter_plot = axis.scatter(survey_df['My_R'],
                             survey_df['My_Python'],
                             facecolors = 'none',
                             edgecolors = 'blue')
plt.savefig('plot-scatter-r-python.pdf',
            bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
            orientation='portrait', papertype=None, format=None,
            transparent=True, pad_inches=0.25, frameon=None)




survey_df_labels = [
    'Personal Preference for Java/Scala/Spark',
    'Personal Preference for Java/Script/HTML/CSS',
    'Personal Preference for Python',
    'Personal Preference for R',
    'Personal Preference for SAS',
    'Professional Java/Scala/Spark',
    'Professional JavaScript/HTML/CSS',
    'Professional Python',
    'Professional R',
    'Professional SAS',
    'Industry Java/Scala/Spark',
    'Industry Java/Script/HTML/CSS',
    'Industry Python',
    'Industry R',
    'Industry SAS'
]



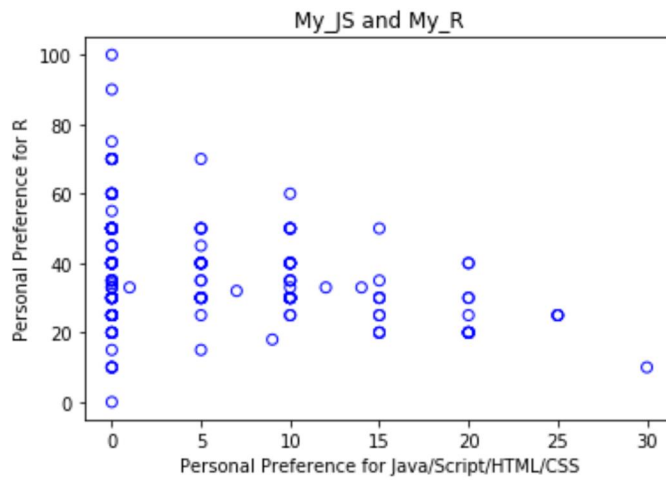
```









```
In [113]: # create a set of scatter plots for personal preferences
for i in range(5):
    for j in range(5):
        if i != j:
            file_title = survey_df.columns[i] + '_and_' + survey_df.columns[j]
            plot_title = survey_df.columns[i] + ' and ' + survey_df.columns[j]
            fig, axis = plt.subplots()
            axis.set_xlabel(survey_df_labels[i])
            axis.set_ylabel(survey_df_labels[j])
            plt.title(plot_title)
            scatter_plot = axis.scatter(survey_df[survey_df.columns[i]],
                                         survey_df[survey_df.columns[j]],
                                         facecolors = 'none',
                                         edgecolors = 'blue')
            plt.savefig(file_title + '.pdf',
                        bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                        orientation='portrait', papertype=None, format=None,
                        transparent=True, pad_inches=0.25, frameon=None)
```












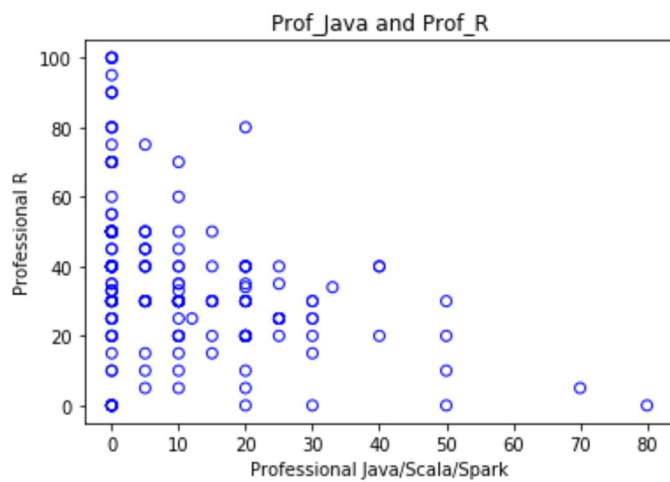

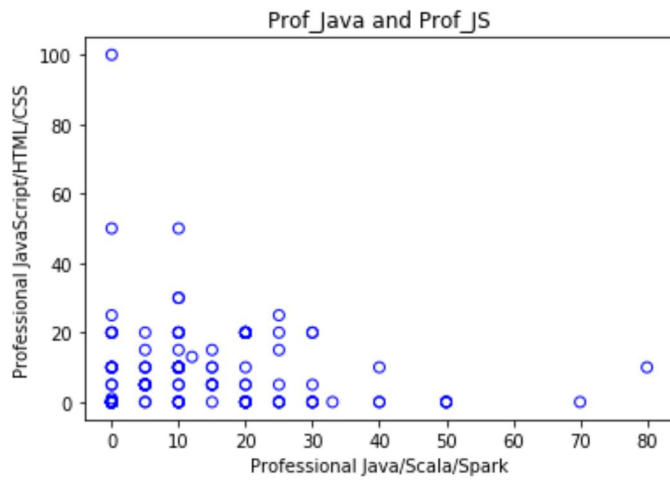


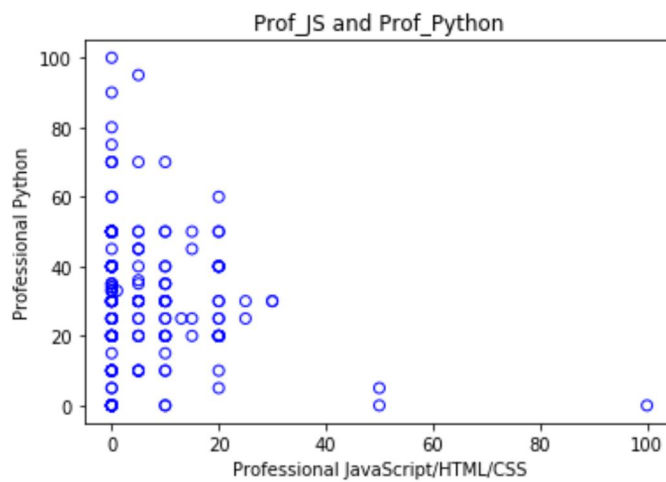






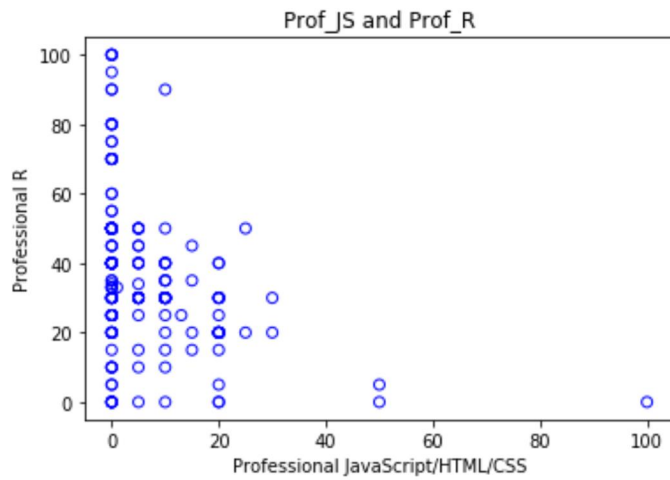










```
In [17]: # create a set of scatter plots for professional preferences
for i in range(5,10):
    for j in range(5,10):
        if i != j:
            file_title = survey_df.columns[i] + '_and_' + survey_df.columns[j]
            plot_title = survey_df.columns[i] + ' and ' + survey_df.columns[j]
            fig, axis = plt.subplots()
            axis.set_xlabel(survey_df_labels[i])
            axis.set_ylabel(survey_df_labels[j])
            plt.title(plot_title)
            scatter_plot = axis.scatter(survey_df[survey_df.columns[i]],
                                         survey_df[survey_df.columns[j]],
                                         facecolors = 'none',
                                         edgecolors = 'blue')
            plt.savefig(file_title + '.pdf',
                        bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                        orientation='portrait', papertype=None, format=None,
                        transparent=True, pad_inches=0.25, frameon=None)
```
































```
In [18]: # create a set of scatter plots for industry preferences
for i in range(11,15):
    for j in range(11,15):
        if i != j:
            file_title = survey_df.columns[i] + '_and_' + survey_df.columns[j]
            plot_title = survey_df.columns[i] + ' and ' + survey_df.columns[j]
            fig, axis = plt.subplots()
            axis.set_xlabel(survey_df_labels[i])
            axis.set_ylabel(survey_df_labels[j])
            plt.title(plot_title)
            scatter_plot = axis.scatter(survey_df[survey_df.columns[i]],
                                         survey_df[survey_df.columns[j]],
                                         facecolors = 'none',
                                         edgecolors = 'blue')
            plt.savefig(file_title + '.pdf',
                        bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                        orientation='portrait', papertype=None, format=None,
                        transparent=True, pad_inches=0.25, frameon=None)
```










```
In [19]: # examine intercorrelations among software preference variables
# with correlation matrix/heat map
corr_chart(df_corr = software_df)
```

<Figure size 432x288 with 0 Axes>



```
In [20]: # descriptive statistics for software preference variables
print('\nDescriptive statistics for survey data -----')
print(software_df.describe())
```

Descriptive statistics for survey data -----

| | My_Java | My_JS | My_Python | My_R | My_SAS | Prof_Java \ |
|-------|------------|------------|------------|------------|------------|-------------|
| count | 207.000000 | 207.000000 | 207.000000 | 207.000000 | 207.000000 | 207.000000 |
| mean | 10.135266 | 4.797101 | 31.304348 | 37.125604 | 16.637681 | 9.251208 |
| std | 11.383477 | 6.757764 | 15.570982 | 14.576003 | 13.626400 | 13.167505 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 0.000000 | 20.000000 | 30.000000 | 5.000000 | 0.000000 |
| 50% | 9.000000 | 0.000000 | 30.000000 | 35.000000 | 15.000000 | 5.000000 |
| 75% | 20.000000 | 10.000000 | 40.000000 | 50.000000 | 25.000000 | 15.000000 |
| max | 70.000000 | 30.000000 | 90.000000 | 100.000000 | 75.000000 | 80.000000 |

| | Prof_JS | Prof_Python | Prof_R | Prof_SAS | Ind_Java \ |
|-------|------------|-------------|------------|------------|------------|
| count | 207.000000 | 207.000000 | 207.000000 | 207.000000 | 207.000000 |
| mean | 5.840580 | 30.028986 | 36.415459 | 18.463768 | 11.942029 |
| std | 10.812555 | 19.144802 | 20.847606 | 18.831841 | 14.706399 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 20.000000 | 25.000000 | 0.000000 | 0.000000 |
| 50% | 0.000000 | 30.000000 | 33.000000 | 15.000000 | 5.000000 |
| 75% | 10.000000 | 40.000000 | 50.000000 | 30.000000 | 20.000000 |
| max | 100.000000 | 100.000000 | 100.000000 | 100.000000 | 70.000000 |

| | Ind_JS | Ind_Python | Ind_R | Ind_SAS |
|-------|------------|------------|------------|------------|
| count | 207.000000 | 207.000000 | 207.000000 | 207.000000 |
| mean | 6.966184 | 29.772947 | 32.434783 | 18.884058 |
| std | 10.030721 | 17.959816 | 15.912209 | 19.137623 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 20.000000 | 22.500000 | 0.000000 |
| 50% | 0.000000 | 30.000000 | 30.000000 | 15.000000 |
| 75% | 10.000000 | 40.000000 | 40.000000 | 30.000000 |
| max | 50.000000 | 95.000000 | 85.000000 | 100.000000 |

```
In [22]: # descriptive statistics for one variable
print('\nDescriptive statistics for courses completed -----')
print(survey_df['Courses_Completed'].describe())
```

Descriptive statistics for courses completed -----

| | |
|-------|------------|
| count | 187.000000 |
| mean | 6.342246 |
| std | 3.170849 |
| min | 1.000000 |
| 25% | 4.000000 |
| 50% | 6.000000 |
| 75% | 9.000000 |
| max | 12.000000 |

Name: Courses_Completed, dtype: float64

```
In [24]: # graduation date counts
grad_date= list(valid_survey_input.Graduate_Date)
from collections import Counter
Counter(grad_date)
```

```
Out[24]: Counter({nan: 3,
                  'Spring 2018': 30,
                  'Fall 2018': 20,
                  'Fall 2017': 14,
                  'Summer 2019': 3,
                  'Summer 2018': 11,
                  'Winter 2019': 11,
                  '2020 or Later': 5,
                  'Winter 2018': 25,
                  'Fall 2019': 5,
                  'Fall 2016': 13,
                  'Summer 2017': 14,
                  'Spring 2019': 9,
                  'Spring 2017': 19,
                  'Winter 2017': 25})
```

```
In [145]: # descriptive statistics for interest in new courses
newcourse_df = survey_df.loc[:, 'Python_Course_Interest':'Systems_Analysis_Course_Interest']

print('\nDescriptive statistics for new courses -----')
print(newcourse_df.describe())
```

```
Descriptive statistics for new courses -----
Python_Course_Interest  Foundations_DE_Course_Interest  \
count                  206.000000                    200.000000
mean                   73.529126                     58.045000
std                    29.835429                     32.588079
min                    0.000000                      0.000000
25%                   53.000000                     29.500000
50%                   82.500000                     60.000000
75%                  100.000000                     89.250000
max                   100.000000                    100.000000

Analytics_App_Course_Interest  Systems_Analysis_Course_Interest
count                        203.000000                    200.000000
mean                        55.201970                     53.630000
std                         34.147954                     33.539493
min                         0.000000                      0.000000
25%                        25.000000                     21.500000
50%                        60.000000                     51.500000
75%                        85.000000                     80.250000
max                        100.000000                    100.000000
```

```
In [28]: # -----
# transformation code added with version v005
# -----
# transformations a la Scikit Learn
# documentation at http://scikit-learn.org/stable/auto\_examples/preprocessing/plot\_all\_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py
#
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
```


```
In [29]: # transformations a la Scikit Learn
# select variable to examine, eliminating missing data codes
X = survey_df['Courses_Completed'].dropna()
```

```
In [30]: # Seaborn provides a convenient way to show the effects of transformations
# on the distribution of values being transformed
# Documentation at https://seaborn.pydata.org/generated/seaborn.distplot.html

unscaled_fig, ax = plt.subplots()
sns.distplot(X).set_title('Unscaled')
unscaled_fig.savefig('Transformation-Unscaled' + '.pdf',
                    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                    orientation='portrait', papertype=None, format=None,
                    transparent=True, pad_inches=0.25, frameon=None)
```

C:\Users\Jimmy\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



```
In [32]: standard_fig, ax = plt.subplots()
sns.distplot(StandardScaler().fit_transform(np.array(X).reshape(-1,1))).set_title('
My_Python StandardScaler')
standard_fig.savefig('Transformation-StandardScaler' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

C:\Users\Jimmy\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



```
In [33]: minmax_fig, ax = plt.subplots()
sns.distplot(MinMaxScaler().fit_transform(np.array(X).reshape(-1,1))).set_title('My
_Python MinMaxScaler')
minmax_fig.savefig('Transformation-MinMaxScaler' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

C:\Users\Jimmy\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



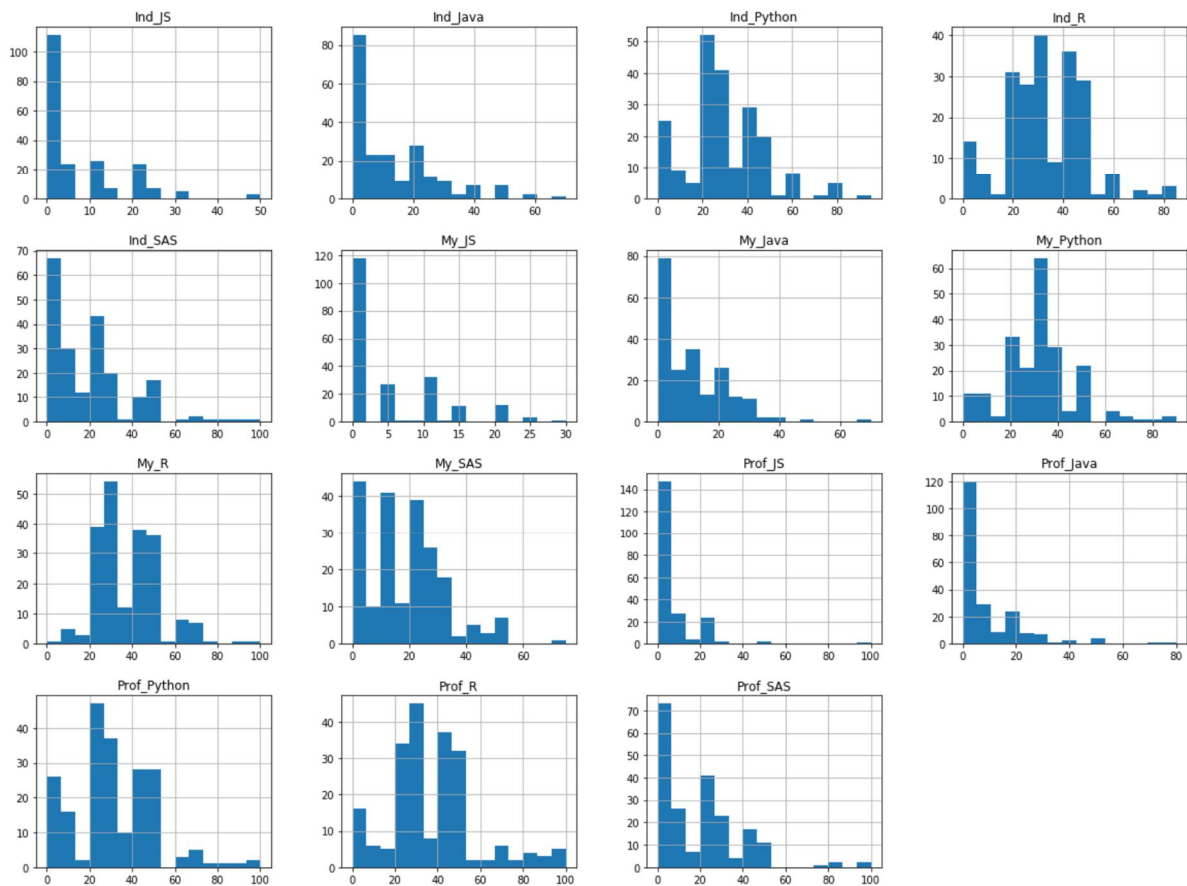
```
In [34]: log_fig, ax = plt.subplots()
sns.distplot(np.log(X)).set_title('My_Python NaturalLog')
log_fig.savefig('Transformation-NaturalLog' + '.pdf',
               bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
               orientation='portrait', papertype=None, format=None,
               transparent=True, pad_inches=0.25, frameon=None)
```

C:\Users\Jimmy\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```




```
In [43]: # histogram of score distribution for each software
software_df.hist(bins = 15, figsize = (20, 15))
plt.show()
```




```
In [41]: survey_df['Courses_Completed'].hist(bins = 25, figsize = (20, 15), color="red")  
plt.show()
```




```
In [57]: # bar plot that shows sum of scores given to software
software_df.sum().plot( kind='bar', figsize = (20, 15) )
plt.title("Sum of Scores by Software", fontsize=20)
plt.xlabel('Software', fontsize=20)
plt.ylabel('Sum of Scores', fontsize=20)

plt.show()
```



```
In [84]: # bar plot that shows mean of scores given to software
software_df.mean().plot( kind='bar', figsize = (20, 15) )
plt.title("Mean of Scores by Software", fontsize=20)
plt.xlabel('Software', fontsize=20)
plt.ylabel('Mean of Scores', fontsize=20)

plt.show()
```



```
In [78]: #Sum of software type
Java_sum= software_df.My_Java.sum() + software_df.Prof_Java.sum()+ software_df.Ind_
Java.sum()
JS_sum= software_df.My_JS.sum() + software_df.Prof_JS.sum()+ software_df.Ind_JS.sum
()
Python_sum= software_df.My_Python.sum() + software_df.Prof_Python.sum()+ software_d
f.Ind_Python.sum()
R_sum= software_df.My_R.sum() + software_df.Prof_R.sum()+ software_df.Ind_R.sum()
SAS_sum= software_df.My_SAS.sum() + software_df.Prof_SAS.sum()+ software_df.Ind_SA
S.sum()


list={'Name': ['Java', 'JS', 'Python', 'R', 'SAS'],
      'Sum': [Java_sum, JS_sum, Python_sum, R_sum, SAS_sum]}
software_type_df= (pd.DataFrame(list))
software_type_df
```

Out[78]:

| | Name | Sum |
|---|--------|-------|
| 0 | Java | 6485 |
| 1 | JS | 3644 |
| 2 | Python | 18859 |
| 3 | R | 21937 |
| 4 | SAS | 11175 |

```
In [87]: #plot of sum by type of software
software_type_df.plot( kind='bar', figsize = (20, 15), legend = False )
plt.title("Sum of Scores by Software Type", fontsize=20)
plt.xlabel('Software Type', fontsize=20)
plt.ylabel('Sum of Scores', fontsize=20)
plt.xticks(np.arange(5), ('Java', 'JS', 'Python', 'R', 'SAS'))


plt.show()
```



```
In [135]: # define subset DataFrame for analysis of software preferences and courses completed
course_comp_df = survey_df.loc[:, 'My_Java': 'Courses_Completed']
cc_software=course_comp_df.groupby("Courses_Completed").sum().loc[:, 'My_Java': 'Ind_SAS']
```

```
In [137]: cc_software.plot( kind='barh', figsize = (10,15), legend = True )
plt.title("Software Preference by Courses Completed ", fontsize=20)
plt.xlabel('Sum of Scores', fontsize=20)
plt.ylabel('Amount of Courses Completed', fontsize=20)


plt.show()
```




```
In [146]: #Analysis of course preference

newcourse_df.sum().plot( kind='bar', figsize = (20, 15) )
plt.title("Sum of Scores by Course", fontsize=20)
plt.xlabel('Course', fontsize=20)
plt.ylabel('Sum of Scores', fontsize=20)

plt.show()
```





```
In [147]: newcourse_df.mean().plot( kind='bar', figsize = (20, 15) )  
plt.title("Mean of Scores by Course", fontsize=20)  
plt.xlabel('Course', fontsize=20)  
plt.ylabel('Mean of Scores', fontsize=20)  
  
plt.show()
```



```
In [148]: cc_preference=course_comp_df.groupby("Courses_Completed").sum().loc[:,
        'Python_Course_Interest':'Systems_Analysis_Course_Interest']
cc_preference.plot( kind='barh', figsize = (10,15),legend = True )
plt.title("Course Preference by Courses Completed ", fontsize=20)
plt.xlabel('Sum of Scores', fontsize=20)
plt.ylabel('Amount of Courses Completed', fontsize=20)

plt.show()
```



In []: